



BCG concluded -0.53 . Ten months of fresh data says -0.51 .

The price-elasticity model BCG built for Evidensia, now re-run on our own data through April 2026. The headline question for pricing: does BCG's conclusion still hold when the model runs on live, growing data we own and validate ourselves? It does — and here is the evidence behind it.



A stable average can hide where the movement sits

The headline says “nothing happened.” Decomposing the drift by service and by revenue weight shows that is almost true — and pinpoints the one place that did move.

**-0.494 →
-0.497**

Median elasticity, 2025 baseline → today. Effectively unchanged.

46% / 48%

Split stronger / weaker — direction is balanced. Net revenue-weighted drift is +0.020: negligible.

1.6%

of rows drift beyond 1.0 — and these concentrate in low-revenue, weak-signal items.

Largest revenue services — stability vs movement

Consult · 1.6 bn		drift +0.167
Imaging · 964 M		+0.003
Surgery · 868 M		+0.000
Internal · 837 M		−0.011
Hospitalisation · 810 M		+0.000
Anaesthesia · 647 M		+0.010
Other · 559 M		−0.011
Dental · 533 M		+0.007
Prescription · 448 M		+0.001

Bar = median |drift| in elasticity (scaled). Of the nine largest services, eight are essentially flat. Consult shows the most code-level movement — but in **both directions** depending on the code, not a single clear shift.

What the numbers say

The raw median moving from **-0.494** to **-0.497** would tempt the conclusion that nothing changed. Looked at closely, that conclusion largely holds: products split almost evenly between slightly stronger (46%) and slightly weaker (48%), and the revenue-weighted net drift is **+0.020 — negligible**.

No large service changed direction clearly. The five biggest by revenue — Consult, Imaging, Surgery, Internal, Hospitalisation, together over 4.7 bn SEK — are essentially flat. The model's reading of how customers respond to price on the **core assortment** is the same today as it was for BCG.

Conclusion

Price sensitivity on the core assortment held still for ten months. The movements that exist are small and sit off to the side — the next page shows they are either low-revenue or noise.

Where movement exists, it is noise or low-revenue — not the core

The products that moved most sort into clear groups. Separating them is what turns a convenient “nothing changed” into a trustworthy one: the strong, revenue-bearing signals stayed put.

NOISE · DISCOUNT

Weak-signal items settled down

The largest raw drift sits in **Healthcare** (+0.97) — but it is only **34 M SEK** and its 2025 values were extreme outliers near the model's bounds. The extra data pulled them back toward sensible levels. This is **weak-signal stabilisation**, not changed customer behaviour.

SBAS0487 -5.668 → -1.041 (1.9 M)

SBCS0124 -8.537 → -6.915 (8.0 M)

So what: nothing to act on — and fresh data **taming** extreme estimates rather than creating new ones is a sign of a well-behaved model.

LOW-REVENUE · WATCH

Two small categories moved cleanly

Accessories became clearly more price-sensitive (-0.19 → -0.50; 92% of items stronger) and **Consumables** uniformly less so (-0.49 → -0.31). These are the only **directionally consistent** shifts — but both are small (26–49 M SEK), so the business impact is limited.

Accessories -0.188 → -0.501 (26 M, 92% stronger)

Consumables -0.486 → -0.305 (49 M, 0% stronger)

So what: worth a look when revisiting pricing in those categories; not material to the top line.

Why this separation is the proof of soundness

A model that drifted randomly would scatter its largest movements across high-revenue services with no pattern. Instead the movement sorts cleanly: **noise-normalisation confined to tiny, weak-signal items**, and **clean directional shifts only in small categories**. The big revenue-bearing services — Consult, Imaging, Surgery, Internal, Hospitalisation — stayed put. That is exactly how a trustworthy elasticity model should behave on fresh data: the signals that matter hold, and the only thing that really moves is the noise getting quieter.

Triangulated conclusion

The fresh-data model is sound: stable where it matters, self-correcting where it was weak, with no large service shifting direction. Prices can be set on these numbers — with a light watch on Accessories and Consumables.

Price Elasticity · Confidential – internal

Movement ranked by drift × revenue (decision impact)

Why these numbers can be trusted

Confidence comes from two things: the sheer volume of data behind the model, and the fact that every step was independently validated — including proving we reproduce BCG exactly before trusting the model on new data.

108,979

product × site × cluster rows in the final blended output

15,128

distinct products carrying a fresh elasticity

7

clinic clusters · 4 years of weekly transactions (2022–2026)

14 / 14

curated validation checks passing across four suites

The validation chain (curated — full detail in the Excel appendix)

Layer	What it proves (in business terms)	Headline	Status
Replication	The full F1–F7 weave reproduces BCG's final elasticity per product, bit-for-bit, on 2025 data	corr 1.000000 · 108,979 rows	PASS
Replication	Cluster, site and bundle models each reproduce BCG's elasticities	3812 / 4673 / 125 groups	PASS
DW extraction	Our own data warehouse reproduces BCG's revenue & volume, per product code	13,223 codes · Σ matches	PASS
Rationality	Every fresh elasticity is negative and within the sensible band — no absurd values	100% negative · 100% in band	PASS
Rationality	Drift vs the 2025 baseline stays within the expected snapshot band	95% drift < 0.5	PASS
Freshness	The core revenue-weighted elasticity is stable: BCG $-0.532 \rightarrow$ today -0.512	net drift $+0.020$	PASS
Provenance	Exactly which inputs are fresh vs frozen — fully transparent, nothing hidden	2 fresh · 3 frozen (documented)	REVIEW*

* REVIEW is by design: three inputs (revenue weights, service routing, bundle branch) remain on BCG's 2025 values to deliver a fresh read now. Their combined effect on decisions is small and each has a defined path to fresh. This is documented, not a defect — the provenance check exists precisely to make it visible.

What is not yet fresh — and the path

Service routing (2025 granularity map) — cheapest to refresh.
Revenue weights (2025 turnover) — moderate rebuild.
Bundle branch — sources just 2.2% of decisions; parked on evidence.

Lift in that order. Until then, the core elasticities — the figures that actually drive pricing — are already fresh.

Method & reservations

Grain: product × site × cluster, blended through a 7-level fallback weave. Elasticity = log-log price coefficient (negative = price-sensitive). Drift = today – 2025, revenue-weighted for decision impact. Weak-signal tail items move most and are discarded before any price decision. Stability over time is a reproduction claim, not a causal price-setting claim. Capital/impact figures are order-of-magnitude.